



MNIST Dataset

The MNIST database of handwritten digits (<http://yann.lecun.com>)

[Data Card](#)[Code \(225\)](#)[Discussion \(0\)](#)[Suggestions \(0\)](#)

About Dataset

Context

MNIST is a subset of a larger set available from NIST (it's copied from <http://yann.lecun.com/exdb/mnist/>)

Content

The MNIST database of handwritten digits has a training set of 60,000 examples, and a test set of 10,000 examples. .

Four files are available:

- train-images-idx3-ubyte.gz: training set images (9912422 bytes)
- train-labels-idx1-ubyte.gz: training set labels (28881 bytes)
- t10k-images-idx3-ubyte.gz: test set images (1648877 bytes)
- t10k-labels-idx1-ubyte.gz: test set labels (4542 bytes)

How to read

See [sample MNIST reader](#)

Usability ⓘ

7.50

License


Data files © Original Authors

Expected update frequency

Not specified

Tags

[Computer Science](#)[Image](#)[Multiclass Classification](#)



Trianing Data

- 60,000 images of handwritten digits from 250 individuals. Among these, 50% are high school students and the other 50% are employees from the Census Bureau.
- All images are composed of 28×28 pixels. Each image also has a corresponding label indicating the digit represented in the image.

Testing Data

- The test data contains 10,000 images of handwritten digits, also sourced from U.S. high school students and Census Bureau employees.
- These images come from a different set of 250 individuals than those whose handwriting was included in the training data.

Why MNIST Dataset?

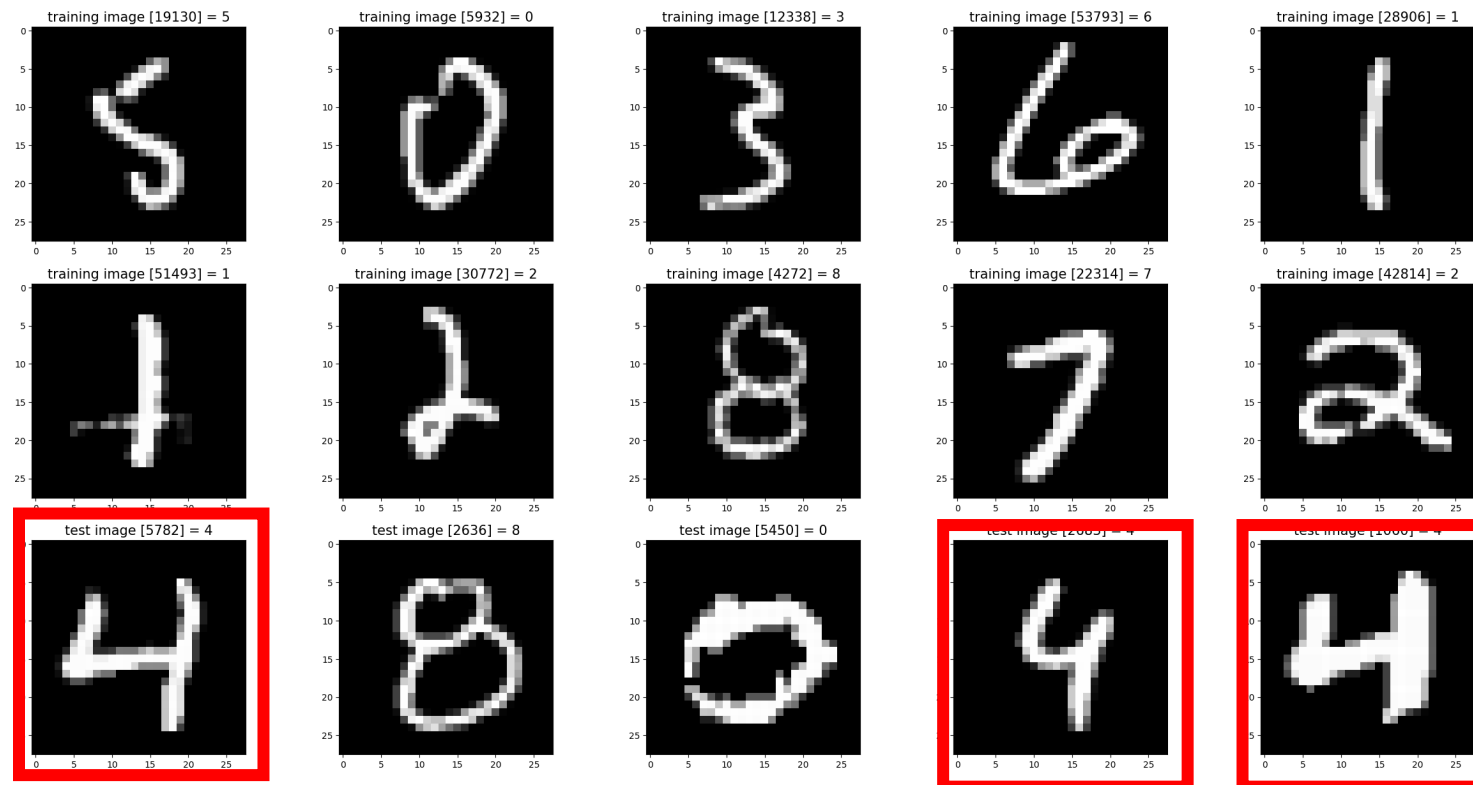
- **Data Availability:** The MNIST dataset, which stands for the Modified National Institute of Standards and Technology dataset, is widely used for training and testing in the field of machine learning. It is a large database of handwritten digits that is publicly available and thus easy to access.
- **Well-defined Problem:** Since the dataset and problem (digit recognition) are well-defined, you can focus more on the modeling aspects rather than spending a lot of time in the data gathering and cleaning stages.

Identify a Supervised Learning Problem

- Why supervised machine learning?
 - MNist is a labeled dataset
 - Room for improvement: top deep learning methods like CNN achieved over 99% accuracy. There is little room for improvement.
 - Faster than deep learning: Compared to CNN, supervised machine learning methods like SVM has shorter run time.

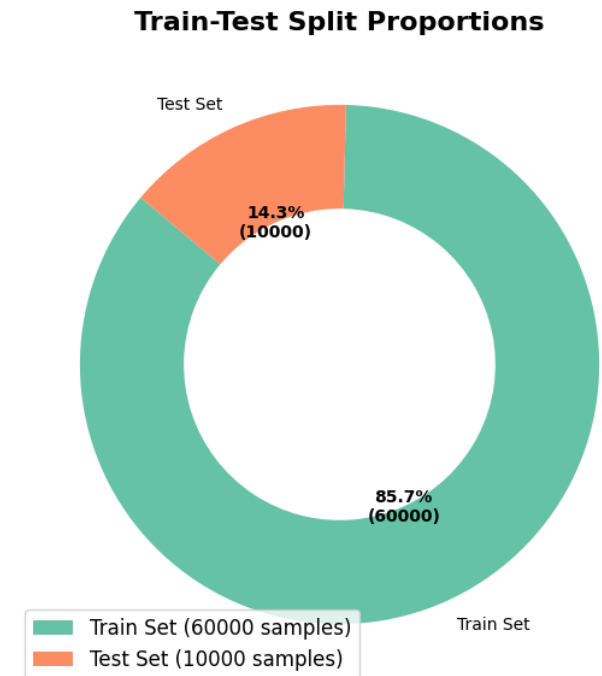
EDA

- I first take a look into the input data
- Most of the handwritings are easy to identify by human eyes
- But I do see a variety of writing styles



EDA

- Each input picture is 224x224x3
 - Most pixels are black and do not contain much information
- Dataset is well balanced among all 10 digits
- Pre-defined test/train split
- At first glance, I believe no transform is required, as MNist is a clean dataset



Models & Results & Analysis: Supervised Method

- Directly apply linear kernel SVM on MNist dataset
 - Full Features: $224 \times 224 \times 3 = 150528$
 - Acc: 91%
 - Run time: 20 mins
 - Session was shut down several times due to excessive memory usage.

Supervised Machine learning: SVM

```
from sklearn import svm
from sklearn import metrics
svm_linear = svm.SVC(kernel='linear')
print("Full Features: ", X_train_tsne.shape)
print(X_val.shape)
y_val_label = np.argmax(y_val, axis=1)
```

Full Features: (4800, 150528)
(1200, 224, 224, 3)

+ Code

+ Markdown

```
# fit
svm_linear.fit(X_train_tsne, y_train_label)
predictions_SVM = svm_linear.predict(X_val.reshape(1200, -1))
correct=0
for i in range(len(predictions_SVM)):
    if(predictions_SVM[i]==y_val_label[i]):
        correct+=1
print(correct/len(predictions_SVM))
```

0.9141666666666667

+ Code

+ Markdown

Models & Results & Analysis:

Supervised Method with PCA Dimension Reduction

- Use PCA to reduce feature size to 20 per image
 - Lower feature size(~5) produces low accuracy (~66%)
- Apply SVM on the reduced features
 - Acc: 91%
 - Run time: 11 s

Supervised Machine learning: SVM with PCA dimension reduction

+ Code + Markdown

```
[21]: from sklearn.decomposition import PCA
pca = PCA(n_components = 20)
pca.fit(X_train_tsne)
X_train_pca = pca.transform(X_train_tsne)
print(X_train_pca.shape)
```

(4800, 20)

```
[22]: from sklearn import svm
svm_linear2 = svm.SVC(kernel='linear')
start = timeit.default_timer()
# fit
svm_linear2.fit(X_train_pca, y_train_label)

print("train_complete")
X_val_pca = pca.transform(X_val.reshape(1200,-1))
print(X_val_pca.shape)
predictions_pca = svm_linear2.predict(X_val_pca)

correct=0
for i in range(len(predictions_pca)):
    if(predictions_pca[i]==y_val_label[i]):
        correct+=1
print(correct/len(predictions_pca))

stop = timeit.default_timer()
print('Time: ', stop - start)
```

train_complete
(1200, 20)
0.9108333333333334
Time: 11.623891282000159

Models & Results & Analysis:

Supervised Method with Hyperparameter Tuning

- Apply grid search on C, gamma
 - Best combination:
 - C=5, gamma=0.001
 - Accuracy 94%

Supervised Machine learning: SVM hyperparameter tuning

```
[23]: from sklearn.model_selection import GridSearchCV
parameters = {'C': [1, 5, 10],
             'gamma': [0.001, 0.01, 1]}
model_SVM_GS = svm.SVC()
grid = GridSearchCV(estimator=model_SVM_GS, param_grid=parameters)
grid.fit(X_train_pca, y_train_label)
print(grid)
# summarize the results of the grid search
print(grid.best_score_)
print(grid.best_estimator_)

GridSearchCV(estimator=SVC(),
              param_grid={'C': [1, 5, 10], 'gamma': [0.001, 0.01, 1]})
0.944375
SVC(C=5, gamma=0.001)
```

```
[24]: predictions_grid = grid.predict(X_val_pca)

correct=0
for i in range(len(predictions_grid)):
    if(predictions_grid[i]==y_val_label[i]):
        correct+=1
print(correct/len(predictions_grid))

0.9483333333333334
```

Models & Results & Analysis:

Unsupervised Method

- Unsupervised methods can also be applied on this dataset
- We use first TSNE to perform dimension reduction
- Followed by GMM clustering
- Ended with label permutation to match with ground truth

```
X_train_tsne=X_train.reshape(4800, -1)
print(X_train_tsne.shape)
tsne = TSNE(n_components=3,
            n_jobs=-1)
digits_proj = tsne.fit_transform(X_train_tsne)
gmm = GaussianMixture(
    n_components=10)
clusters = gmm.fit_predict(digits_proj)
```

Models & Results & Analysis: Unsupervised Method Performance

- The accuracy is 83.6%
- Run time ~ 2 min
- Increasing TSNE component does not improve accuracy

```
stop = timeit.default_timer()

print('Time: ', stop - start)
```

```
(4800, 224, 224, 3)
(4800, 150528)
Time: 157.55008647299974
```

```
voter=np.zeros(shape=(10,10))
y_train_label=np.argmax(y_train,axis=1)
for i in range(len(clusters)):
    #print(clusters[i],y_train_label[i])
    voter[clusters[i]][y_train_label[i]]+=1
print(voter)
trans=np.argmax(voter,axis=1)
correct=0
for i in range(len(clusters)):
    if y_train_label[i]==trans[clusters[i]]:
        correct+=1
print(correct/len(clusters))
```

```
[[ 1.  1.  6. 448.  0. 29.  0.  0. 33.  7.]
 [ 0.  3.  9.  4. 13.  0.  0. 470.  1. 159.]
 [ 1.  0.  3.  2. 10.  7. 479.  0.  3.  1.]
 [476.  0.  1.  1.  0.  1.  3.  0.  7.  5.]
 [ 0. 206.  8.  3.  4.  1.  1.  4.  5.  0.]
 [ 0.  2. 429.  1.  1.  1.  0.  4.  2.  0.]
 [ 0.  1.  2.  7.  0.  1.  0.  0. 347.  0.]
 [ 5.  1.  3.  8.  0. 366.  4.  1. 32.  0.]
 [ 0.  1.  1.  3. 472. 10.  1. 25.  8. 299.]
 [ 1. 322.  4.  4.  4.  1.  1.  5.  2.  2.]]
0.8364583333333333
```

Models & Results & Analysis:

Unsupervised Method using less training data


- Although the accuracy is lower, unsupervised methods is less sensitive to insufficient training data
- We experiment using only half of the training set for training
- Accuracy drops from 83.6% to 83.2%
- While run time drops from 2 min to less than 1 min

```
(2400, 150528)
Time: 54.03585478500008
[[ 1.  2.  0.  4. 243.  7.  2. 10.  3. 174.]
 [ 0.  0.  5.  7.  0. 170.  4.  1.  2.  0.]
 [232.  0.  0.  0.  0.  0.  2.  0.  1.  1.]
 [ 0.  1.  5. 223.  0. 23.  0.  0. 10.  3.]
 [ 1. 147.  3.  0.  0.  1.  0.  5.  1.  0.]
 [ 0.  2.  9.  1. 13.  0.  0. 220.  0. 41.]
 [ 0.  1. 218.  6.  0.  0.  0.  3.  5.  0.]
 [ 0.  0.  1.  2.  6.  6. 241.  0.  4.  0.]
 [ 0. 110.  3.  2.  2.  0.  0.  3.  2.  0.]
 [ 0.  0.  0.  7.  0.  4.  0.  0. 194.  0.]]
0.8325
```

Models & Results & Analysis:


Deep-Learning Method

- Since CNN is the king of 2D image processing, I decided to include it in the comparison
- Since 224x224 is the input size of the famous VGG16, we use the pre-trained model provided in Keras to solve the problem.
- Below is the architecture of VGG16. The schematic consists of hierarchical CNN feature extractor and fully-connected classifier.



Models & Results & Analysis: Deep-Learning Method Performance

- Achieved 99.5% accuracy
- Accuracy is almost 90% at the beginning of training, thanks to the pretrained weight.
- Run time is ~15min



Conclusion

- Raw input, even cleaned, may still need transform like dimension reduction for the following reasons
 - Shorter run time
 - Lower hardware requirement (smaller memory usage)
- Deep learning methods outperforms traditional methods, at the price of higher computing cost
- Hyperparameter selection is crucial regarding model performance

```

import numpy as np
import pandas as pd

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

/kaggle/input/mnist-dataset/train-images.idx3-ubyte
/kaggle/input/mnist-dataset/t10k-labels.idx1-ubyte
/kaggle/input/mnist-dataset/t10k-images.idx3-ubyte
/kaggle/input/mnist-dataset/train-labels.idx1-ubyte
/kaggle/input/mnist-dataset/t10k-labels-idx1-ubyte/t10k-labels-idx1-
ubyte
/kaggle/input/mnist-dataset/t10k-images-idx3-ubyte/t10k-images-idx3-
ubyte
/kaggle/input/mnist-dataset/train-labels-idx1-ubyte/train-labels-idx1-
ubyte
/kaggle/input/mnist-dataset/train-images-idx3-ubyte/train-images-idx3-
ubyte

```

Libraries to be used

```

import tensorflow as tf
from keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense, Dropout
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split

```

Load Data

```

(train_images, train_labels), (test_images, test_labels) =
mnist.load_data()
train_images, test_images = train_images / 255.0, test_images / 255.0
# normalising the dimensions of images to increase the performance of
cnn

# Load the MNIST dataset
(train_images, train_labels), (test_images, test_labels) =
mnist.load_data()

# Normalize the images
train_images, test_images = train_images / 255.0, test_images / 255.0

```

```

# Use only 60% of the dataset
train_images, train_labels = train_images[:int(0.1 *
len(train_images))], train_labels[:int(0.1 * len(train_labels))]

# Convert images to 3-channel (RGB) and resize to 224x224 for VGG16
def preprocess_images(images):
    images = tf.expand_dims(images, axis=-1) # Add channel dimension
    images = tf.image.grayscale_to_rgb(images) # Convert to RGB
    images = tf.image.resize(images, [224, 224]) # Resize to 224x224
    return images.numpy()

train_images = preprocess_images(train_images)
test_images = preprocess_images(test_images)

# Convert labels to one-hot encoding
train_labels = to_categorical(train_labels, num_classes=10)
test_labels = to_categorical(test_labels, num_classes=10)

# Split train data into train and validation sets (80-20 split of 60%
dataset)
X_train, X_val, y_train, y_val = train_test_split(train_images,
train_labels, test_size=0.2, random_state=42)

random_index = np.random.randint(0, len(train_images))
random_image = train_images[random_index]
random_label = train_labels[random_index]

```

EDA

```

temp=np.argmax(train_labels, axis=1)
d={}
for i in range(10):
    d[i]=0
for q in temp:
    d[q]+=1
print(d)
#balanced dataset

{0: 592, 1: 671, 2: 581, 3: 608, 4: 623, 5: 514, 6: 608, 7: 651, 8:
551, 9: 601}

for i in range(20):
    random_index = np.random.randint(0, len(train_images))
    random_image = train_images[random_index]
    random_label = train_labels[random_index]

    plt.imshow(random_image, cmap='gray') # Use 'gray' for MNIST
images

```

```
plt.axis('off')  
plt.title(np.argmax(random_label)) # Convert one-hot label back  
to a class index  
plt.show()
```

3



1



4



8



8



2



8



7



0



4



9



2



8



4




1



4



1



8



0



1



```
random_image.shape
(224, 224, 3)

labels = str(list(range(0,10)))
labels
'[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]'

num_samples = 12
random_indices = np.random.choice(train_images.shape[0], num_samples,
replace=False)
sample_images = train_images[random_indices]
sample_labels = train_labels[random_indices]

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

Deep learning: VGG16 Network

```
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Flatten, Dense, Dropout
vgg_base = VGG16(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))
```

```

# Freeze the VGG16 base layers
for layer in vgg_base.layers:
    layer.trainable = False

# Add custom layers on top of VGG16
x = Flatten()(vgg_base.output)
x = Dense(128, activation='relu')(x)
x = Dropout(0.5)(x)
output = Dense(10, activation='softmax')(x) # 10 classes for MNIST

# Create final model
model = Model(inputs=vgg_base.input, outputs=output)
# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])

# Train the model without using ImageDataGenerator
##change to true if train CNN
if False:
    history = model.fit(
        X_train, y_train, # Directly use the training data
        validation_data=(X_val, y_val), # Directly use the validation
data
        epochs=10,
        batch_size=32 # Set batch size here
    )

```

Unsupervised method: TSNE+GMM

```

from sklearn.manifold import Isomap, TSNE
from sklearn.decomposition import PCA
# Clustering
from sklearn.mixture import GaussianMixture
import timeit

print(X_train.shape)
X_train_tsne=X_train.reshape(4800,-1)
print(X_train_tsne.shape)
y_train_label=np.argmax(y_train,axis=1)
y_val_label=np.argmax(y_val,axis=1)

(4800, 224, 224, 3)
(4800, 150528)

if True:#change to true if you want to run GMM
    start = timeit.default_timer()

```

```

tsne = TSNE(n_components=3,
            n_jobs=-1)
digits_proj = tsne.fit_transform(X_train_tsne)
gmm = GaussianMixture(
    n_components=10)
clusters = gmm.fit_predict(digits_proj)
#Your statements here

stop = timeit.default_timer()

print('Time: ', stop - start)
voter=np.zeros(shape=(10,10))

for i in range(len(clusters)):
    #print(clusters[i],y_train_label[i])
    voter[clusters[i]][y_train_label[i]]+=1
print(voter)
trans=np.argmax(voter,axis=1)
correct=0
for i in range(len(clusters)):
    if y_train_label[i]==trans[clusters[i]]:
        correct+=1
print(correct/len(clusters))

```

```

Time: 157.3312752239999
[[ 0.  0.  1.  0. 267.  0.  0. 26.  2. 264.]
 [ 0. 219.  7.  3.  4.  1.  1.  4.  4.  0.]
 [476.  0.  0.  1.  0.  1.  3.  0.  4.  5.]
 [ 1. 308.  2.  3.  4.  1.  1.  5.  3.  1.]
 [ 0.  2.  9.  3.  1.  0.  0. 441.  1. 10.]
 [ 5.  1.  5. 13.  3. 247. 12.  1. 358.  0.]
 [ 0.  3. 430.  2.  1.  1.  0.  3.  2.  0.]
 [ 1.  1.  4.  2.  7.  4. 471.  0.  2.  1.]
 [ 1.  1.  7. 450.  0. 155.  1.  0. 59.  7.]
 [ 0.  2.  1.  4. 217.  7.  0. 29.  5. 185.]]
0.7577083333333333

```

Unsupervised method: reduce training data

```

start = timeit.default_timer()

X_train_tsne2=X_train_tsne[:2400][:]
y_train_label2=np.argmax(y_train,axis=1)[:2400]
print(X_train_tsne2.shape)
tsne2 = TSNE(n_components=3,
            n_jobs=-1)
digits_proj2 = tsne2.fit_transform(X_train_tsne2)
gmm2 = GaussianMixture(

```



```

    n_components=10)
clusters2 = gmm2.fit_predict(digits_proj2)
#Your statements here

stop = timeit.default_timer()

print('Time: ', stop - start)

voter2=np.zeros(shape=(10,10))

for i in range(len(clusters2)):
    #print(clusters[i],y_train_label[i])
    voter2[clusters2[i]][y_train_label2[i]]+=1
print(voter2)
trans2=np.argmax(voter2,axis=1)
correct2=0
for i in range(len(clusters2)):
    if y_train_label2[i]==trans2[clusters2[i]]:
        correct2+=1
print(correct2/len(clusters2))

(2400, 150528)
Time: 56.688194474000284
[[ 1. 147.  3.  0.  0.  1.  0.  5.  1.  0.]
 [ 0.  1.  5. 223.  0. 28.  0.  0.  8.  3.]
 [ 0.  0.  1.  2.  6.  7. 241.  0.  4.  0.]
 [ 1.  2.  1.  4. 244.  7.  2. 13.  3. 182.]
 [ 0.  1. 219.  6.  0.  0.  0.  3.  5.  0.]
 [ 0. 110.  3.  2.  2.  0.  0.  3.  2.  0.]
 [ 0.  0.  0.  7.  0.  4.  0.  0. 196.  0.]
 [ 0.  2.  8.  1. 12.  0.  0. 217.  0.  33.]
 [ 0.  0.  4.  7.  0. 164.  4.  1.  2.  0.]
 [232.  0.  0.  0.  0.  0.  2.  0.  1.  1.]]
0.8304166666666667

```

Supervised Machine learning: SVM

```

from sklearn import svm
from sklearn import metrics
svm_linear = svm.SVC(kernel='linear')
print("Full Features: ",X_train_tsne.shape)
print(X_val.shape)
y_val_label=np.argmax(y_val,axis=1)

```

```

Full Features: (4800, 150528)
(1200, 224, 224, 3)

```

```
# fit
if False: #change to True if you want to apply SVM on raw data
    svm_linear.fit(X_train_tsne, y_train_label)
    predictions_SVM = svm_linear.predict(X_val.reshape(1200,-1))
    correct=0
    for i in range(len(predictions_SVM)):
        if(predictions_SVM[i]==y_val_label[i]):
            correct+=1
    print(correct/len(predictions_SVM))
```

Supervised Machine learning: SVM with PCA dimension reduction

```
from sklearn.decomposition import PCA
pca = PCA(n_components = 20)
pca.fit(X_train_tsne)
X_train_pca = pca.transform(X_train_tsne)
print(X_train_pca.shape)

(4800, 20)

from sklearn import svm
svm_linear2 = svm.SVC(kernel='linear')
start = timeit.default_timer()
# fit
svm_linear2.fit(X_train_pca, y_train_label)

print("train_complete")
X_val_pca = pca.transform(X_val.reshape(1200,-1))
print(X_val_pca.shape)
predictions_pca = svm_linear2.predict(X_val_pca)

correct=0
for i in range(len(predictions_pca)):
    if(predictions_pca[i]==y_val_label[i]):
        correct+=1
print(correct/len(predictions_pca))

stop = timeit.default_timer()
print('Time: ', stop - start)

train_complete
(1200, 20)
0.9083333333333333
Time: 13.015603397999712
```

Supervised Machine learning: SVM

hyperparameter tuning

```
from sklearn.model_selection import GridSearchCV
parameters = {'C': [1, 5, 10],
              'gamma': [0.001, 0.01, 1]}
model_SVM_GS = svm.SVC()
grid = GridSearchCV(estimator=model_SVM_GS, param_grid=parameters)
grid.fit(X_train_pca, y_train_label)
print(grid)
# summarize the results of the grid search
print(grid.best_score_)
print(grid.best_estimator_)

GridSearchCV(estimator=SVC(),
              param_grid={'C': [1, 5, 10], 'gamma': [0.001, 0.01, 1]})
0.9441666666666666
SVC(C=5, gamma=0.001)

predictions_grid = grid.predict(X_val_pca)

correct=0
for i in range(len(predictions_grid)):
    if(predictions_grid[i]==y_val_label[i]):
        correct+=1
print(correct/len(predictions_grid))

0.9483333333333334
```

Deep learning: other analysis

```
from tensorflow.keras.optimizers import Adagrad

# Modeling
model.compile(optimizer='Adagrad', loss='categorical_crossentropy',
metrics=['accuracy'])

model.summary()

Model: "functional_3"
```

Layer (type)	Output Shape
Param #	

0	input_layer_1 (InputLayer)	(None, 224, 224, 3)	
	block1_conv1 (Conv2D)	(None, 224, 224, 64)	
1,792			
	block1_conv2 (Conv2D)	(None, 224, 224, 64)	
36,928			
	block1_pool (MaxPooling2D)	(None, 112, 112, 64)	
0			
	block2_conv1 (Conv2D)	(None, 112, 112, 128)	
73,856			
	block2_conv2 (Conv2D)	(None, 112, 112, 128)	
147,584			
	block2_pool (MaxPooling2D)	(None, 56, 56, 128)	
0			
	block3_conv1 (Conv2D)	(None, 56, 56, 256)	
295,168			
	block3_conv2 (Conv2D)	(None, 56, 56, 256)	
590,080			
	block3_conv3 (Conv2D)	(None, 56, 56, 256)	
590,080			
	block3_pool (MaxPooling2D)	(None, 28, 28, 256)	
0			
	block4_conv1 (Conv2D)	(None, 28, 28, 512)	
1,180,160			
	block4_conv2 (Conv2D)	(None, 28, 28, 512)	

2,359,808				
	block4_conv3 (Conv2D)		(None, 28, 28, 512)	
2,359,808				
	block4_pool (MaxPooling2D)		(None, 14, 14, 512)	
0				
	block5_conv1 (Conv2D)		(None, 14, 14, 512)	
2,359,808				
	block5_conv2 (Conv2D)		(None, 14, 14, 512)	
2,359,808				
	block5_conv3 (Conv2D)		(None, 14, 14, 512)	
2,359,808				
	block5_pool (MaxPooling2D)		(None, 7, 7, 512)	
0				
	flatten_1 (Flatten)		(None, 25088)	
0				
	dense_2 (Dense)		(None, 128)	
3,211,392				
	dropout_1 (Dropout)		(None, 128)	
0				
	dense_3 (Dense)		(None, 10)	
1,290				

Total params: 17,927,370 (68.39 MB)

Trainable params: 3,212,682 (12.26 MB)

Non-trainable params: 14,714,688 (56.13 MB)

```
loss, accuracy = model.evaluate(X_train, y_train)
print(f"The model accuracy is : {accuracy} \n the model loss : {loss}")
```

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1733920451.204469 102 service.cc:145] XLA service 0x7dc46000cf40 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1733920451.204534 102 service.cc:153] StreamExecutor device (0): Tesla T4, Compute Capability 7.5

I0000 00:00:1733920451.204540 102 service.cc:153] StreamExecutor device (1): Tesla T4, Compute Capability 7.5

1/150 ————— 36:33 15s/step - accuracy: 0.0625 - loss: 2.5141

I0000 00:00:1733920465.297473 102 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

150/150 ————— 33s 124ms/step - accuracy: 0.0936 - loss: 2.3835

The model accuracy is : 0.10104166716337204
the model loss : 2.377922773361206

```
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label = 'Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```


NameError Traceback (most recent call last)

Cell In[52], line 1

----> 1 plt.plot(history.history['accuracy'], label='Training Accuracy')

2 plt.plot(history.history['val_accuracy'], label = 'Validation Accuracy')

3 plt.xlabel("Epochs")

NameError: name 'history' is not defined

```
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel("Epoche")
plt.ylabel("Loss")
```

```

plt.legend()
plt.show()

import numpy as np

# Predict the class probabilities for the test dataset
predictions = model.predict(X_val)

# Convert probabilities to class labels
predicted_labels = np.argmax(predictions, axis=1)
true_labels = np.argmax(y_val, axis=1) # Convert one-hot to class
indices

from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

# Generate the confusion matrix
true_labels = np.argmax(y_val, axis=1) # Convert one-hot encoded to
class indices
predicted_labels = np.argmax(predictions, axis=1) # Already
converting probabilities to class indices
correct=0
for i in range(len(true_labels)):
    if true_labels[i]==predicted_labels[i]:
        correct+=1
print(correct/len(true_labels))
from sklearn.metrics import confusion_matrix

# Generate confusion matrix
conf_matrix = confusion_matrix(true_labels, predicted_labels)

# Plot the confusion matrix
plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
xticklabels=range(10), yticklabels=range(10))
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()

```